

Dirichlet Process

Il-Chul Moon
Dept. of Industrial and Systems Engineering
KAIST

icmoon@kaist.ac.kr

DEFINITION OF DIRICHLET PROCESS

Detour: Gaussian Mixture Model

- Let's assume that the data points are drawn from a mixture distribution of multiple multivariate Gaussian distributions

- $$P(x) = \sum_{k=1}^K P(z_k)P(x|z) = \sum_{k=1}^K \pi_k N(x|\mu_k, \Sigma_k)$$

- How to model such mixture?

- Mixing coefficient, or Selection variable: z_k

- The selection is stochastic which follows the multinomial distribution

- $$z_k \in \{0,1\}, \sum_k z_k = 1, P(z_k = 1) = \pi_k, \sum_{k=1}^K \pi_k = 1, 0 \leq \pi_k \leq 1$$

- $$P(Z) = \prod_{k=1}^K \pi_k^{z_k}$$

- Mixture component

- $$P(X|z_k = 1) = N(x|\mu_k, \Sigma_k) \rightarrow P(X|Z) = \prod_{k=1}^K N(x|\mu_k, \Sigma_k)^{z_k}$$



- This is the marginalized probability. How about conditional?

- $$\gamma(z_{nk}) \equiv p(z_k = 1|x_n) = \frac{P(z_k=1)P(x|z_k = 1)}{\sum_{j=1}^K P(z_j=1)P(x|z_j = 1)}$$

$$= \frac{\pi_k N(x|\mu_k, \Sigma_k)}{\sum_{j=1}^K \pi_j N(x|\mu_j, \Sigma_j)}$$

- Log likelihood of the entire dataset is

- $$\ln P(X|\pi, \mu, \Sigma) = \sum_{n=1}^N \ln \{ \sum_{k=1}^K \pi_k N(x|\mu_k, \Sigma_k) \}$$



Detour: Dirichlet Distribution



- Generative Process**

- $\theta_i \sim \text{Dir}(\alpha), i \in \{1, \dots, M\}, \varphi_k \sim \text{Dir}(\beta), k \in \{1, \dots, K\}$
- $z_{i,l} \sim \text{Mult}(\theta_i), i \in \{1, \dots, M\}, l \in \{1, \dots, N\}, w_{i,l} \sim \text{Mult}(\varphi_{z_{i,l}}), i \in \{1, \dots, M\}, l \in \{1, \dots, N\}$


- Dirichlet Distribution**

- $P(x_1, \dots, x_K | \alpha_1, \dots, \alpha_K) = \frac{\Gamma(\sum_{i=1}^K \alpha_i)}{\prod_{i=1}^K \Gamma(\alpha_i)} x_i^{\alpha_i - 1}$


- $x_1, \dots, x_{K-1} > 0$
- $x_1 + \dots + x_{K-1} < 1$
- $x_K = 1 - x_1 - \dots - x_{K-1}$
- $\alpha_i > 0$




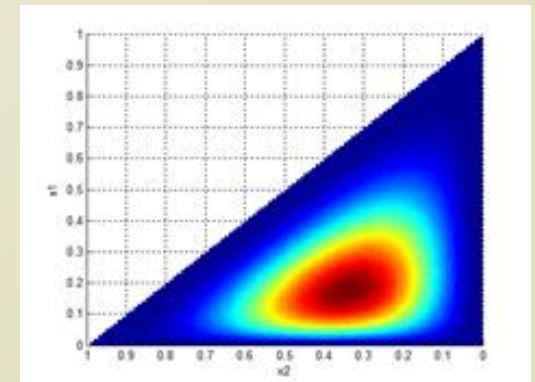
$$[\alpha_1, \alpha_2, \alpha_3] = [1, 1, 1]$$



$$[\alpha_1, \alpha_2, \alpha_3] = [2, 2, 2]$$



$$[\alpha_1, \alpha_2, \alpha_3] = [2, 3, 4]$$



Multinomial-Dirichlet Conjugate Relation

- Multinomial distribution
 - N independently and identically distributed instances, $N = \sum_i c_i$
 - c_i is the number of occurrences of the i-th choice
 - $P(D|\theta) = \frac{N!}{\prod_i c_i!} \prod_i \theta_i^{c_i}$
- Dirichlet distribution
 - $P(\theta|\alpha) = \frac{1}{B(\alpha)} \prod_i \theta_i^{\alpha_i-1}$
- Bayesian Posterior
 - $P(\theta|D, \alpha) \propto P(D|\theta)P(\theta|\alpha) = \frac{N!}{\prod_i c_i!} \prod_i \theta_i^{c_i} \frac{1}{B(\alpha)} \prod_i \theta_i^{\alpha_i-1} = \frac{N!}{B(\alpha) \prod_i c_i!} \prod_i \theta_i^{\alpha_i+c_i-1} \propto \prod_i \theta_i^{\alpha_i+c_i-1}$
 - $P(\theta|D, \alpha) = \frac{1}{B(\alpha+c)} \prod_i \theta_i^{\alpha_i+c_i-1}$
 - Coming back to the Dirichlet distribution : Conjugate Prior
 - The likelihood of the Dirichlet distribution is the conjugate prior of the multinomial distribution
- Dirichlet distribution with D as a single observation with i-th choice
 - $\theta|\alpha \sim \text{Dir}(\alpha_1, \dots, \alpha_i, \dots, \alpha_K)$
 - $\theta|\alpha, D \sim \text{Dir}(\alpha_1, \dots, \alpha_i + 1, \dots, \alpha_K)$

Dirichlet Process


- Dirichlet process, $G|\alpha, H \sim DP(\alpha, H)$
 - $(G(A_1), \dots, G(A_r))|\alpha, H \sim \text{Dir}(\alpha H(A_1), \dots, \alpha H(A_r))$
 - $A_1 \cap \dots \cap A_r = \emptyset, A_1 \cup \dots \cup A_r = \Theta$
 - Properties

$$E[G(A)] = H(A)$$

$$V[G(A)] = \frac{H(A)(1 - H(A))}{\alpha + 1}$$

- H : Base distribution
 - α : Concentration parameter, strength parameter (strength of prior)
- Posterior distribution given a dataset of $\theta_1 \dots \theta_n$
 - $\text{Posterior} \propto \text{Likelihood} \times \text{Prior}$
 - Multinomial-Dirichlet conjugate relationship
 - The posterior becomes the Dirichlet distribution, again, adjusted to reflect the likelihood
 - $(G(A_1), \dots, G(A_r))|\theta_1 \dots \theta_n, \alpha, H \sim \text{Dir}(\alpha H(A_1) + n_1, \dots, \alpha H(A_r) + n_r)$
 - $n_k = |\{\theta_i | \theta_i \in A_k, 1 \leq i \leq n\}|$

$$G|\theta_1 \dots \theta_n, \alpha, H \sim DP\left(\alpha + n, \frac{\alpha}{\alpha + n} H + \frac{n}{\alpha + n} \frac{\sum_{i=1}^n \delta_{\theta_i}}{n}\right)$$



$\text{Dir}(2,3,4)$

Sampling from Dirichlet Process

- Dirichlet process
 - $(G(A_1), \dots, G(A_r)) | \alpha, H \sim \text{Dir}(\alpha H(A_1), \dots, \alpha H(A_r))$
 - $G | \theta_1 \dots \theta_n, \alpha, H \sim \text{DP} \left(\alpha + n, \frac{\alpha}{\alpha + n} H + \frac{n}{\alpha + n} \frac{\sum_{i=1}^n \delta_{\theta_i}}{n} \right)$
- Definition is done, but how to realize the definition?
 - How to draw an instance, or a distribution, G , from the Dirichlet process?
 - How to draw an instance, θ_i , from the distribution, G ?
- Multiple generation *schemes*, or *construction*, exist
 - From the definition of Dirichlet process to the sample from the Dirichlet process
 - Stick Breaking Scheme
 - Polya Urn Scheme
 - Chinese Restaurant Process Scheme

Stick-Breaking Construction

- Imagine that we create a probability mass function on infinite choices

- $k = 1, 2, \dots, \infty$
- $v_k | \alpha \sim \text{Beta}(1, \alpha)$
- $\beta_k = v_k \prod_{l=1}^{k-1} (1 - v_l)$

- Common notation is

- $\beta \sim \text{GEM}(\alpha)$

- We were constructing a distribution for the Dirichlet process

- $G | \alpha, H \sim \text{DP}(\alpha, H)$
 - $\beta \sim \text{GEM}(\alpha)$
 - $G = \sum_{k=1}^{\infty} \beta_k \delta_{\theta_k}$
 - $\theta_k | H \sim H$

- θ_k chooses a n -th broken stick, and the stick length is the prob.
- We know the existence of the infinite-th stick length.

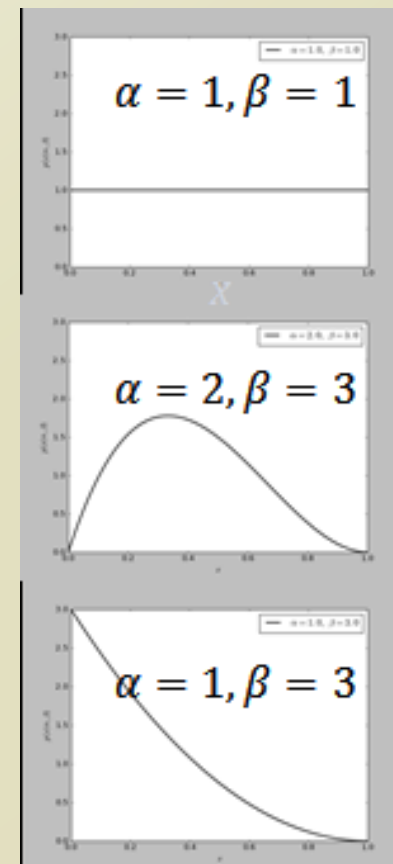
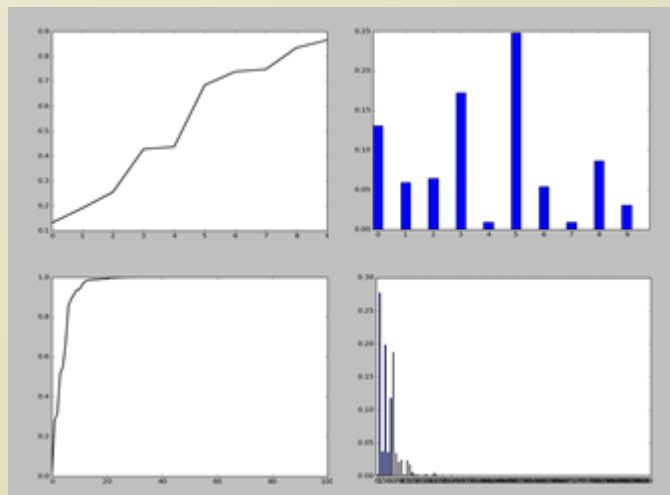

- Exponential growth in CDF

→ Discount the growth


→ Pitman-Yor Process

Close to Power law dist.


Useful for language models...



Polya Urn Scheme



$\alpha = 4$



$\alpha = 10$

- Dirichlet process
 - $G|\theta_1 \dots \theta_n, \alpha, H \sim DP\left(\alpha + n, \frac{\alpha}{\alpha+n} H + \frac{n}{\alpha+n} \frac{\sum_{i=1}^n \delta_{\theta_i}}{n}\right)$
 - $G|\alpha, H \sim DP(\alpha, H)$
 - $(G(A_1), \dots, G(A_r))|\alpha, H \sim \text{Dir}(\alpha H(A_1), \dots, \alpha H(A_r))$
 - $E[G(A)] = H(A)$
 - $\theta_n|\theta_1 \dots \theta_{n-1}, \alpha, H \sim DP\left(\alpha + n - 1, \frac{\alpha}{\alpha+n-1} H + \frac{n-1}{\alpha+n-1} \frac{\sum_{i=1}^{n-1} \delta_{\theta_i}}{n-1}\right)$
 - $E[\theta_n|\theta_1 \dots \theta_{n-1}, \alpha, H] \sim \frac{\alpha}{\alpha+n-1} H + \frac{\sum_{i=1}^{n-1} \delta_{\theta_i}}{\alpha+n-1} \sim \frac{\alpha}{\alpha+n-1} H + \frac{\sum_{k=1}^K N_k \delta_{\theta_k}}{\alpha+n-1}$, N_k : the number of k-th choice occurrences
 - This enables sampling an observation from the Dirichlet process without constructing $G|\alpha, H \sim DP(\alpha, H)$
 - Stick-breaking (distribution) *construction* vs. Polya Urn *sampling* from distribution
- Polya Urn Scheme
 - Create an empty urn
 - Do
 - toss = Coin toss from $[0, \alpha + n - 1]$
 - If $0 \leq \text{toss} < \alpha$
 - Add a ball to the urn by paining the ball as a sample from $\theta_n \sim H$
 - If $\alpha \leq \text{toss} < \alpha + n - 1$
 - Pick a ball from the urn
 - Return the ball and a new ball with the same color to the urn

Chinese Restaurant Process

- Dirichlet process

- $G|\theta_1 \dots \theta_n, \alpha, H \sim DP\left(\alpha + n, \frac{\alpha}{\alpha+n} H + \frac{n}{\alpha+n} \frac{\sum_{i=1}^n \delta_{\theta_i}}{n}\right)$
- $E[\theta_n | \theta_1 \dots \theta_{n-1}, \alpha, H]$

$$\sim \frac{\alpha}{\alpha + n - 1} H + \frac{\sum_{i=1}^{n-1} \delta_{\theta_i}}{\alpha + n - 1}$$

$$\sim \frac{\alpha}{\alpha + n - 1} H + \frac{\sum_{k=1}^K N_k \delta_{\theta_k}}{\alpha + n - 1}$$

N_k : the number of k-th choice occurrences

- $P(\theta_n | \theta_1 \dots \theta_{n-1}, \alpha) = \begin{cases} \frac{N_k}{\alpha + n - 1} \\ \frac{\alpha}{\alpha + n - 1} \end{cases}$


- Chinese restaurant process

- Assume Infinite number of tables in a restaurant
- First customer sits at the first table
- Loop for Customer N sits at:

- 1) Table k with $P(\theta_n | \theta_1 \dots \theta_{n-1}, \alpha) = \frac{N_k}{\alpha + n - 1}$
- 2) A new table $k+1$ with $P(\theta_n | \theta_1 \dots \theta_{n-1}, \alpha) = \frac{\alpha}{\alpha + n - 1}$



- Properties of Chinese restaurant process

- Clustering formation
- Rich-get-richer property
- No fixed number of clusters with a fixed number of instances
- Almost identical to Polya Urn Scheme



Detour: Random Process

- Random process, a.k.a. stochastic process, is
 - An infinite indexed collection of random variables, $\{X(t)|t \in T\}$
 - Index parameter : t
 - Can be time, space....
 - A function, $X(t, \omega)$, where $t \in T$ and $\omega \in \Omega$
 - Outcome of the underlying random experiment : ω
 - Fixed $t \rightarrow X(t, \omega)$ is a random variable over Ω
 - Fixed $\omega \rightarrow X(t, \omega)$ is a deterministic function of t , a sample function
- Example of random process
 - Gaussian process
 - Fixed t , a random variable following a Gaussian distribution
 - Fixed ω , a deterministic curve of t
 - Dirichlet process
 - Fixed t , a random variable following a Dirichlet distribution
 - Fixed ω , a deterministic placement over clusters



de Finetti's Theorem


- Exchangeability
 - A joint probability distribution is exchangeable if it is invariant to permutation
 - Given a permutation of S
 - $P(x_1, x_2, \dots, x_N) = P(x_{S(1)}, x_{S(2)}, \dots, x_{S(N)})$
- (De Finetti, 1935) If (x_1, x_2, \dots) are infinitely exchangeable, then the joint probability $P(x_1, x_2, \dots, x_N)$ has a representation as a mixture

$$P(x_1, x_2, \dots, x_N) = \int \left(\prod_{i=1}^N P(x_i | \theta) \right) dP(\theta) = \int P(\theta) \left(\prod_{i=1}^N P(x_i | \theta) \right) d\theta$$

For some random variable θ


- Independent and identically distributed \rightarrow Exchangeable
- Exchangeable \rightarrow IID : No. A counter example is the Polya urn sampling
- Chinese restaurant process is an exchangeable process
 - No proof in this scope
 - Why is exchangeability important?
 - Enables a simple derivation of Gibbs sampler for the inference
 - We remove the instance of the next Gibbs sampling from the existing cluster assignment

Detour: Concept of Gibbs Sampling



- Each step involves **replacing** the value of one of the variables by a value drawn from the distribution of that variable conditioned on the values of the remaining variables
- Repeated either by cycling through the variables in some particular order or by choosing the variable to be updated at each step at random from some distribution
- Example


1. Full joint probability : $p(z_1, z_2, z_3)$
2. Sample $z_1 \sim p(z_1 | z_2^{(\tau)}, z_3^{(\tau)})$
→ Obtain a value $z_1^{(\tau+1)}$
3. Sample $z_2 \sim p(z_2 | z_1^{(\tau+1)}, z_3^{(\tau)})$
→ Obtain a value $z_2^{(\tau+1)}$
4. Sample $z_3 \sim p(z_3 | z_1^{(\tau+1)}, z_2^{(\tau+1)})$
→ Obtain a value $z_3^{(\tau+1)}$




DIRICHLET PROCESS MIXTURE MODEL

Detour: Gaussian Mixture Model

- Let's assume that the data points are drawn from a mixture distribution of multiple multivariate Gaussian distributions
 - $P(x) = \sum_{k=1}^K P(z_k)P(x|z) = \sum_{k=1}^K \pi_k N(x|\mu_k, \Sigma_k)$
 - How to model such mixture?
 - Mixing coefficient, or Selection variable: z_k
 - The selection is stochastic which follows the multinomial distribution
 - $z_k \in \{0,1\}, \sum_k z_k = 1, P(z_k = 1) = \pi_k, \sum_{k=1}^K \pi_k = 1, 0 \leq \pi_k \leq 1$
 - $P(Z) = \prod_{k=1}^K \pi_k^{z_k}$
 - Mixture component
 - $P(X|z_k = 1) = N(x|\mu_k, \Sigma_k) \rightarrow P(X|Z) = \prod_{k=1}^K N(x|\mu_k, \Sigma_k)^{z_k}$
 - This is the marginalized probability. How about conditional?
 - $$\gamma(z_{nk}) \equiv p(z_k = 1|x_n) = \frac{P(z_k=1)P(x|z_k = 1)}{\sum_{j=1}^K P(z_j=1)P(x|z_j = 1)}$$
$$= \frac{\pi_k N(x|\mu_k, \Sigma_k)}{\sum_{j=1}^K \pi_j N(x|\mu_j, \Sigma_j)}$$
 - Log likelihood of the entire dataset is
 - $\ln P(X|\pi, \mu, \Sigma) = \sum_{n=1}^N \ln \{ \sum_{k=1}^K \pi_k N(x|\mu_k, \Sigma_k) \}$



Dirichlet Process Mixture Model



- Common usage of Dirichlet process : Prior on parameters of a mixture model

- Like $P(z_k = 1) = \pi_k$

- $z_k \in \{0,1\}, \sum_k z_k = 1, P(z_k = 1) = \pi_k, \sum_{k=1}^K \pi_k = 1, 0 \leq \pi_k \leq 1$

- Indicator representation of GMM with infinite K

- $\beta | \gamma \sim GEM(\gamma), \theta_k | H, \lambda \sim H(\lambda), z_i | \beta \sim \beta, x_i | \{\theta_k\}_{k=1}^{\infty}, z_i \sim F(\theta_{z_i})$

- $\beta \sim GEM(\alpha) \rightarrow k = 1, 2, \dots, \infty, v_k | \alpha \sim Beta(1, \alpha), \beta_k = v_k \prod_{l=1}^{k-1} (1 - v_l)$

- Alternative representation of GMM with infinite K

- $G_0 | H, \gamma \sim DP(\gamma, H), \theta'_i | G_0 \sim G_0, x_i | \theta'_i \sim F(\theta'_i)$

- $\theta_n | \theta_1 \dots \theta_{n-1}, \gamma, H \sim DP\left(\gamma + n - 1, \frac{\gamma}{\gamma + n - 1} H + \frac{n-1}{\gamma + n - 1} \frac{\sum_{i=1}^{n-1} \delta_{\theta_i}}{n-1}\right)$

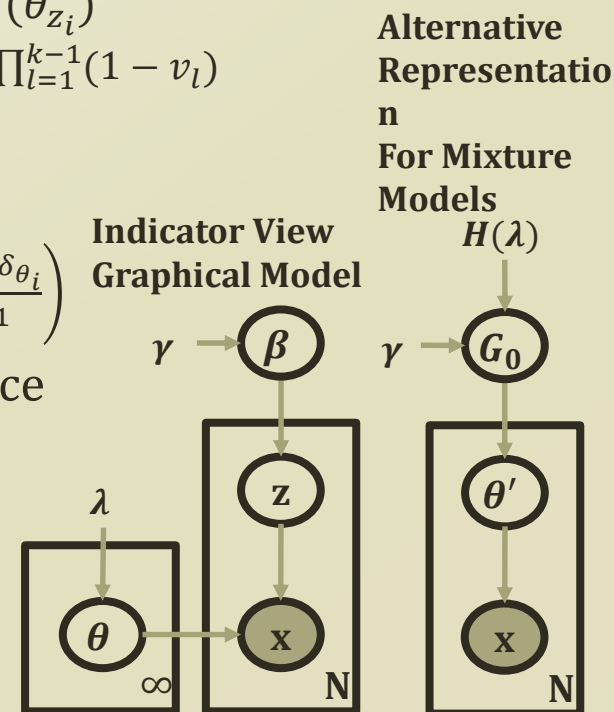
- Continuously updating the assignment of an instance

- Learning concept

- de Finetti's theorem + Chinese restaurant process
 - + Gibbs Sampling

- Each assignment

- Surely updates the parameter of each cluster
 - May create a new cluster



Implementation Details of DPMM

- Online update of the component parameter

- $G_0 | H, \gamma \sim DP(\gamma, H), \theta'_i | G_0 \sim G_0, x_i | \theta'_i \sim F(\theta'_i)$

- $\theta_n | \theta_1 \dots \theta_{n-1}, \gamma, H \sim DP\left(\gamma + n - 1, \frac{\gamma}{\gamma + n - 1} H + \frac{n-1}{\gamma + n - 1} \frac{\sum_{i=1}^{n-1} \delta_{\theta_i}}{n-1}\right), P(\theta_n | \theta_1 \dots \theta_{n-1}, \alpha) = \begin{cases} \frac{N_k}{\alpha + n - 1} \\ \frac{\alpha}{\alpha + n - 1} \end{cases}$

- $F(x_i | \theta'_i) = N(x_i | \mu_{\theta'_i}, \Sigma_{\theta'_i})$

- $\mu_{\theta'_i}$ and $\Sigma_{\theta'_i}$ are the component parameters given that the component follows the Gaussian distribution

- DPMM**

- Initial table assignments**

- While sampling iterations**


- While each data instance in the dataset**

- Remove the instance from the assignment
 - Calculate the prior : $\theta_n | \theta_1 \dots \theta_{n-1}, \gamma, H \sim DP$
 - Calculate the likelihood : $N(x_i | \mu_{\theta'_i}, \Sigma_{\theta'_i})$
 - Calculate the posterior
 - Sample the cluster assignment from the posterior
 - Update the component parameter

- Truncated Dirichlet process mixture model


- Finish the sampling of stick-breaking with the limit on the number of atoms

- Same as limiting the table numbers

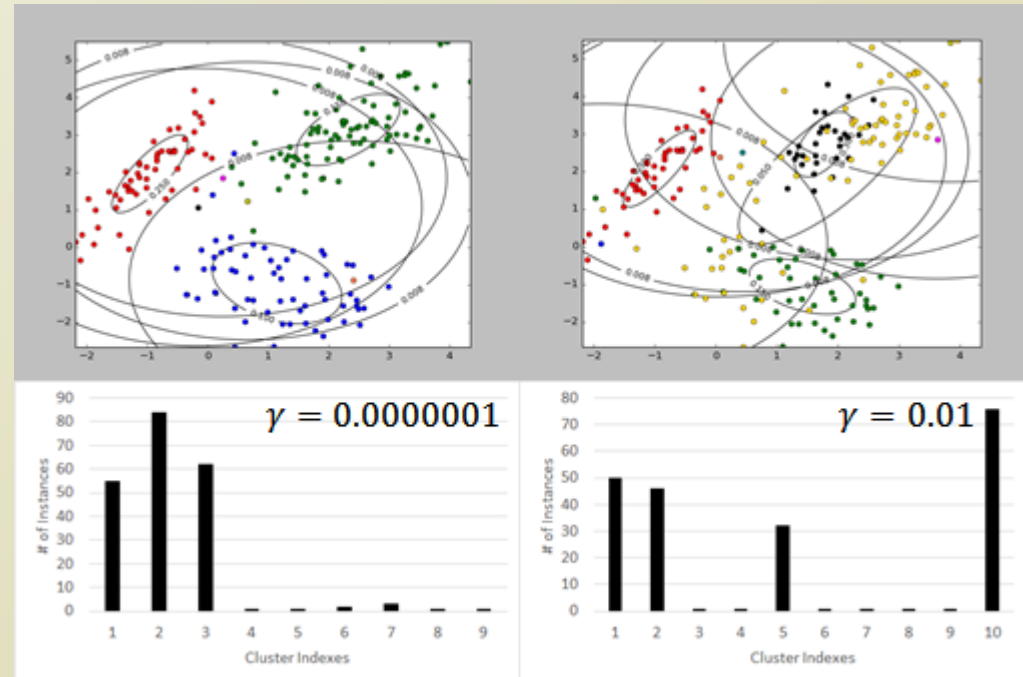


DPMM Sampling Process

- The Sampling process produces the different clustering results per iterations
 - γ can determine the sensitivity of the cluster generation
 - $\theta_n | \theta_1 \dots \theta_{n-1}, \gamma, H \sim DP \left(\gamma + n - 1, \frac{\gamma}{\gamma + n - 1} H + \frac{n-1}{\gamma + n - 1} \frac{\sum_{i=1}^{n-1} \delta_{\theta_i}}{n-1} \right)$




Synthesized True Dataset






HIERARCHICAL DIRICHLET PROCESS

Problem of Separate Prior


- Datasets are often structured
 - LDA : Corpus-Document structure
 - Hierarchical structure
- Finite dimension of clusters
 - Choice is finite, and the atoms will overlap
 - Infinite model might have zero overlap in atoms
 - Smooth continuous distribution of the base distribution
 - Need to enforce sharing the atoms
- Clustering result is different from one branch to another
 - Need to share the same dimension of clusters
 - How to correlate θ_1 and θ_2



Solution of Atom Sharing




- Hierarchical structure of Dirichlet processes
 - H : the continuous base distribution
 - G_0 : a draw from $G_0 \sim DP(H, \gamma)$
 - G_i : a draw from $G_i | G_0 \sim DP(G_0, \alpha_0)$
- Here, G_0 is a discrete distribution
 - so G_i will only sample from the atoms of G_0



Stick Breaking Construction

- A hierarchical Dirichlet process with a corpus with D documents
 - Can be applied to domains other than texts
 - $G_0 \sim \text{DP}(H, \gamma)$
 - $G_i | G_0 \sim \text{DP}(G_0, \alpha_0)$
- Stick breaking (*prior distribution*) construction of HDP
 - $G_0 = \sum_{k=1}^{\infty} \beta_k \delta_{\phi_k}$
 $\phi_k \sim H$ is shared
 - $\phi_k \sim H$
 - $\beta_k = \beta'_k \prod_{l=1}^{k-1} (1 - \beta'_l)$
 - $\beta'_k | \gamma \sim \text{Beta}(1, \gamma)$
 - $G_i = \sum_{k=1}^{\infty} \pi_{ik} \delta_{\phi_k}$
 - $\pi_{ik} = \pi'_{ik} \prod_{l=1}^{k-1} (1 - \pi'_{il})$
 - $\pi'_{ik} | \gamma \sim \text{Beta}(\alpha_0 \beta_k, \alpha_0 (1 - \sum_{i=1}^k \beta_i))$



Chinese Restaurant Franchise

- $G_0 \sim \text{DP}(H, \gamma)$
- $G_i | G_0 \sim \text{DP}(G_0, \alpha_0)$
 - $\theta_{in} \sim G_i$: a θ_{in} 's seating on a ψ_{it} table of each restaurant
 - $\psi_{it} \sim G_0$: a ψ_{it} 's table serves a ϕ_k menu of the franchise

